

Diff finding a needle in a haystack

Take a look at the two files `haystack.txt` and `needle_haystack.txt`. It is difficult to manually find small errors in large output files.

Finding Differences

To find any differences in the files we can simply run the `diff` command

```
$ diff haystack.txt needle_haystack.txt
122c122
< hay
---
> needle
```

This output shows that the files differ at line 122 of each file. The difference is that at line 122 `haystack.txt` has `hay` and `haystack_needle.txt` has `needle`. We know this because of the directions of the angle brackets.

Adding Some Color

I personally like color coding in my tools. `diff`, like some other utilities has a `--color` option. There are 3 settings always, never, and auto. While you might always want color you probably actually want auto. auto will always show colors in the terminal and won't show colors when outputting to a file.

You can make a `diff` alias to always turn on colors for you by editing your shell's rc file, which is probably `~/.bashrc`

```
alias diff="diff --color=auto"
```

Side-by-Side Comparison

If we want to compare the two files side-by-side, which may be more intuitive for some, we can use the `-y` option (there is also a long form `--side-by-side`).

```
$ diff -y haystack.txt needle_haystack.txt
hay                hay
hay                hay
hay                hay
hay                hay
hay                hay
hay                hay
hay                hay
hay                hay
hay                hay
hay                hay
...
```

This generates a full side-by-side comparison of the two files, but maybe you only want to see the lines that are different. to do that you can also add the `--suppress-common-lines` option.

```
$ diff -y --suppress-common-lines haystack.txt needle_haystack.txt
hay                                     | needle
```

Getting Some Context

Maybe you want to see the lines around the line that is different to know why it may be different. We could just open the files in an editor and look for the difference there, but there's an easy way to view the lines around the differences. The `-c` option shows the lines from each file separately whereas the `-u` option shows the lines in a unified view

```
$ diff -c haystack.txt needle_haystack.txt
*** haystack.txt      2019-02-01 10:41:28.722803889 -0500
--- needle_haystack.txt 2019-02-01 10:41:19.236118886 -0500
*****
*** 119,125 ****
    hay
    hay
    weird data
!  hay
    hay
    hay
    hay
--- 119,125 ----
    hay
    hay
    weird data
!  needle
    hay
    hay
    hay
```

Here we can see the files from line 119-125. The lines that are different are marked with a `!`. You can also pass an argument to either `-c` or `-u` to specify the number of lines to show

```
$ diff -u5 haystack.txt needle_haystack.txt
--- haystack.txt      2019-02-01 10:41:28.722803889 -0500
+++ needle_haystack.txt 2019-02-01 10:41:19.236118886 -0500
@@ -117,11 +117,11 @@
    hay
    hay
    hay
    hay
    weird data
-hay
+needle
    hay
    hay
    hay
    hay
    hay
```

Here we can see from line 117 in each file +11 lines.

